

Parth Bhoiwala

Udacity AI Nanodegree (May 2017)
AI Planning Search - Heuristic Analysis

The goal of this project was to define a group of problems in classical PDDL (Planning Domain Definition Language) for the air cargo domain discussed in the lectures. A planning graph was used along with many search methods and heuristic functions to implement the plan. There were total of three air cargo problems and below are the problem's initial state and goal.

Problem 1:

```
Init(At(C1, SFO) ^ At(C2, JFK)
    ^ At(P1, SFO) ^ At(P2, JFK)
    ^ Cargo(C1) ^ Cargo(C2)
    ^ Plane(P1) ^ Plane(P2)
    ^ Airport(JFK) ^ Airport(SFO))
Goal(At(C1, JFK) ^ At(C2, SFO))
```

Problem 2:

```
Init(At(C1, SFO) ^ At(C2, JFK) ^ At(C3, ATL)
    ^ At(P1, SFO) ^ At(P2, JFK) ^ At(P3, ATL)
    ^ Cargo(C1) ^ Cargo(C2) ^ Cargo(C3)
    ^ Plane(P1) ^ Plane(P2) ^ Plane(P3)
    ^ Airport(JFK) ^ Airport(SFO) ^ Airport(ATL))
Goal(At(C1, JFK) ^ At(C2, SFO) ^ At(C3, SFO))
```

Problem 3:

```
Init(At(C1, SFO) ^ At(C2, JFK) ^ At(C3, ATL) ^ At(C4, ORD)
    ^ At(P1, SFO) ^ At(P2, JFK)
    ^ Cargo(C1) ^ Cargo(C2) ^ Cargo(C3) ^ Cargo(C4)
    ^ Plane(P1) ^ Plane(P2)
    ^ Airport(JFK) ^ Airport(SFO) ^ Airport(ATL) ^ Airport(ORD))
Goal(At(C1, JFK) ^ At(C3, JFK) ^ At(C2, SFO) ^ At(C4, SFO))
```

For each of the problems mentioned above, the plan search was computed and the results are recorded below.

Problem 1 was run for all the search functions and Table 1 below shows the results.

Table 1: Air Cargo Problem 1 Results

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
Breadth First	6	0.054	43	56	180
Breadth First Tree	6	1.069	1458	1459	5960
Depth First Graph	12	0.009	12	13	48
Depth Limited	50	0.101	101	271	414
Uniform Cost	6	0.041	55	57	224
Recursive Best First (h1)	6	3.091	4229	4230	17029
Greedy Best First Graph (h1)	6	0.006	7	9	28

As it can be seen, the optimal solutions include the ones with Breadth First, Breadth First Tree, Uniform Cost, Recursive Best First (h1) and Greedy Best First Graph (h1) search. The reason why these are considered optimal is because their plan length is the lowest at 6. Other search methods have plan lengths of 12 and 50 which is a lot of unnecessary work. Out of the five optimal solutions, the best solution would be Greedy Best First Graph search with h1. Not only it provides a plan length of 6 but it's execution time was the fastest at 0.006 seconds and only expanded 7 nodes. The solution provided by Greedy BFG search is shown below:

```
Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, JFK)
Fly(P2, JFK, SFO)
Unload(C1, P1, JFK)
Unload(C2, P2, SFO)
```

Problem 2 was run for Breadth First, Depth First, Uniform and Greedy Best First searches only. The other search functions were terminated after running for more than 10 minutes.

Table 2: Air Cargo Problem 2 Results

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
Breadth First	9	14.26	3401	4672	31049
Depth First Graph	346	1.523	350	351	3142
Uniform Cost	9	12.84	4761	4763	43206
Greedy Best First Graph (h1)	9	1.425	550	552	4950

For problem 2, the optimal solutions include Breadth First, Uniform Cost and Greedy BFG search as they all yield a plan with a plan length of 9. DFG Search yielded a plan with 346 which is too inefficient. Out of the 3 optimal searches, Greedy BFG search wins again this time with the fastest execution time of 1.425 seconds. Below is the solution provided by Greedy BFG search:

```
Load(C1, P1, SFO)
Load(C2, P2, JFK)
Load(C3, P3, ATL)
Fly(P1, SFO, JFK)
Fly(P2, JFK, SFO)
Fly(P3, ATL, SFO)
Unload(C3, P3, SFO)
Unload(C2, P2, SFO)
Unload(C1, P1, JFK)
```

Problem 3 was also run for Breadth First, Depth First, Uniform and Greedy Best First searches only. The other search functions were terminated after running for more than 10 minutes.

Table 3: Air Cargo Problem 3 Results

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
Breadth First	12	99.32	14491	17947	128184
Depth First Graph	1878	19.06	1948	1949	16253
Uniform Cost	12	52.99	17783	17785	155920
Greedy Best First Graph (h1)	22	12.05	4031	4033	35794

For problem 3, Breadth First and Uniform Cost were optimal because they had smallest plan length of 12. Depth First Graph is too inefficient because it's plan length was 1878. Now, Greedy Best First Graph yielded a plan length of 22 which could or could not be considered optimal. This is tricky because a little more analysis is required. Breadth First and Uniform Cost, which yielded a plan length of 12 had execution times of 99.32 and 52.99 seconds respectively. On the other hand, Greedy BFG search, which yielded a plan length of 22 had an execution time of only 12.05 seconds.

So the question to ask here is *"What's more important? Plan length or Execution time?"*.

Because Greedy BFG search wins by providing the fastest execution time of 12.05 seconds which is 4 times as fast as Uniform cost and 9 times faster than Breadth first search. The catch is that Greedy BFG search costs 10 extra steps and if that doesn't matter, then it wins otherwise Uniform Cost wins. Below are the solutions for Uniform Cost and Greedy BFG search.

Uniform Cost

```
Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, ATL)
Load(C3, P1, ATL)
Fly(P2, JFK, ORD)
Load(C4, P2, ORD)
Fly(P1, ATL, JFK)
Fly(P2, ORD, SFO)
Unload(C4, P2, SFO)
Unload(C3, P1, JFK)
Unload(C2, P2, SFO)
Unload(C1, P1, JFK)
```

Greedy BFG with h1

```
Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, ORD)
Load(C4, P1, ORD)
Fly(P2, JFK, ATL)
Load(C3, P2, ATL)
Fly(P2, ATL, SFO)
Fly(P1, ORD, ATL)
Unload(C4, P1, ATL)
Fly(P1, ATL, ORD)
Fly(P2, SFO, ATL)
Load(C4, P2, ATL)
Fly(P2, ATL, ORD)
Unload(C3, P2, ORD)
Load(C3, P1, ORD)
Fly(P2, ORD, SFO)
Fly(P1, ORD, JFK)
Unload(C3, P1, JFK)
Unload(C1, P1, JFK)
Fly(P1, JFK, ORD)
Unload(C4, P2, SFO)
Unload(C2, P2, SFO)
```

After running these uninformed search functions, informed search functions were run on those three problems.

Below is the result for running problem 1 on A* with h1, A* with ignore_preconditions, and A* with level_sum.

Table 4: Problem 1 - Informed Search Results (Air Cargo Problem)

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
A* with h1	6	0.036	55	57	224
A* with ignore_preconditions	6	0.041	41	43	170
A* with level_sum	6	0.979	11	13	50

The results in the above table 5 indicate something interesting. It can be concluded that all three search algorithms were optimal because they all yielded solutions with a plan length of 6. However, A* with h1 executed the problem in 0.036 seconds which is significantly faster compared to A* with level_sum which took 0.979 seconds to run. So clearly, A* with h1 wins. Although, another interesting thing to note is that A* with level_sum expanded only 11 nodes during its execution time. Below is the plan suggested by A* with h1:

```

Load(C1, P1, SFO)
Load(C2, P2, JFK)
Fly(P1, SFO, JFK)
Fly(P2, JFK, SFO)
Unload(C1, P1, JFK)
Unload(C2, P2, SFO)

```

Same searches were run again for Problem 2 and below are the results for that.

Table 5: Problem 2 - Informed Search Results (Air Cargo Problem)

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
A* with h1	9	13.03	4761	4763	43206
A* with ignore_preconditions	9	4.684	1450	1452	13303
A* with level_sum	9	197.1	86	88	841

For problem 2, the results are a little different from the previous one. The plan length for all three searches are 9 which makes all of them optimal. However, unlike the previous one in which A* with h1 provided the fastest execution, it was almost 3 times slower than A* with ignore_preconditions. A* with h1 executed in 13.03 seconds while A* with ignore_preconditions executed in only 4.684 seconds making it the winner. However, one interesting thing to note here is that the number of nodes expanded is also much larger for A* with ignore_preconditions at 1450 while only 4761 for A* with h1. Below is the plan suggested by A* with ignore_preconditions:

```

Load(C3, P3, ATL)
Fly(P3, ATL, SFO)
Unload(C3, P3, SFO)
Load(C2, P2, JFK)
Fly(P2, JFK, SFO)
Unload(C2, P2, SFO)
Load(C1, P1, SFO)
Fly(P1, SFO, JFK)
Unload(C1, P1, JFK)

```

Problem 3 was run for only A* with h1 and A* with ignore_preconditions. It was terminated for A* with level_sum after 10 minutes.

Table 5: Problem 3 - Informed Search Results (Air Cargo Problem)

Search	Plan Length	Time(s)	Expansions	Goal Tests	New Nodes
A* with h1	12	54.44	17783	17785	155920
A* with ignore_preconditions	12	17.33	5003	5005	44586

As it can be seen above, both A* with h1 and A* with ignore_preconditions are optimal with a plan length of 12; however, A* with ignore_preconditions wins easily by providing an execution time of 17.33 seconds. Below is the plan:

```

Load(C2, P2, JFK)
Fly(P2, JFK, ORD)
Load(C4, P2, ORD)
Fly(P2, ORD, SFO)
Unload(C4, P2, SFO)
Load(C1, P1, SFO)
Fly(P1, SFO, ATL)
Load(C3, P1, ATL)
Fly(P1, ATL, JFK)
Unload(C3, P1, JFK)
Unload(C2, P2, SFO)
Unload(C1, P1, JFK)

```

In conclusion, it is fair to say that Greedy Breadth First Graph search with H1 was the most optimal because not only it founded the solution in least amount of time but it also calculated the smallest possible plans in terms of length (with a small exception in problem 3) discussed above. For the informed searches that were run, it was noticed that A* with ignore_preconditions was optimal and the best choice as it was also the fastest. In the first problem, it was slightly slower than A* with h1, however, it was fastest in average case. It is proved that having the best possible search doesn't mean it's the best solution. Having a good heuristic function is vital and can create a significant impact on the overall implementation.